Report: Hook Generation For Music Composition

Julian Napier, Jonas Veit

Universität Heidelberg

March 5, 2025

Contents

1   Goal Of The Project

Popular music nowadays is very dependant on catchy and repetetive melodies, so called "hooks". These 8 bar long melodies what make a song recognizable and consequently sell well. What if there was a way for songwriters to press a button and a generate a suggestion for such a hook? The suggestion wouldn't have to be that fancy; writing a chord progession and transposing it into any desired key isn't a hard task for any skilled music composer. Additionally, it would only serve as inspiration: if the songwriter were to dislike any part of the generated sequence, he could change it to his liking, thus expressing his own creativity and artistic freedom.

Well, that is exactly what we tried to accomplish with our project. As an input to a decoder transformer model in numpy, we feed hooks that have gone through a strict filtering process, to make them as uniform as possible. This means all data is 8 bars, monophonic, in the key of C major or A minor and has a time signature of 4/4 at 120 bpm. In addition, we adapted the tokenization, attention and inference as well as possible to our task. This way, we hoped the transformer would pick up relevant melodic and rhythmic patterns, the generated output being of similar shape.

2   Data Preprocessing And Collection

2.1   The Lakh MIDI Dataset

The Lakh MIDI dataset is a collection of 176,581 multi-track MIDI files. Its goal was to facilitate large-scale music information retrieval, both symbolic and audio content-based (Raffel 2016).

For our project, we primarily used the clean MIDI subset of the dataset, which contains 34,518 files, matched to already known songs. This automatically implies better data quality. Another benefit is that the genre distribution is known; here the top 25 genres, ranked by occurences on the website *Lakh Clean Midi Dataset Analysis* (2024):

1. Pop (2013)
2. Rock (875)
3. Pop rock (538)
4. Alternative/Indie (458)
5. Jazz (336)
6. R&B/Soul (307)
7. Soft rock (264)
8. Hard rock (259)
9. Country (210)
10. Classical (198)
11. Alternative rock (180)
12. Schlager/Volksmusik (176)
13. Rock and roll (170)
14. R&B (159)
15. Folk rock (158)
16. Progressive rock (151)
17. Dance/Electronic (144)
18. Children's music (125)
19. Soul music (119)
20. Synth-pop (112)
21. Hip-hop (107)
22. New age (101)
23. Disco (100)
24. Classic rock (97)
25. Metal (87)

It is evident that the subset covers a wide spectrum of genres. This enables our model to be exposed to a greater variety of data and generalize to common patterns in multiple genres. As can be seen, songs from more popular genres (Pop, Rock) are also recognized more often. This leads to more training data from these genres, which is big upside for the generation of hooks, since they depend on recognizability and catchiness.

2.2 Prefiltering The MIDI Files

```python
def check_time_sig(pretty_midi_stream):
    # Get the list of time signature changes
    time_signature_changes = pretty_midi_stream.time_signature_changes
    tempos = get_tempo(pretty_midi_stream)
    # Check if there are no tempo/time signature changes
    if len(time_signature_changes) == 1 and len(tempos) == 1:
        time_signature = time_signature_changes[0]
        # Check if the time signature is 2/4 or 4/4
        if time_signature.denominator == 4:
            if time_signature.numerator == 4:
                return True
            if time_signature.numerator == 2: #Treat it like 4/4
                pretty_midi_stream.time_signature_changes[0].numerator
                    = 4
                return True
    return False
```

In order for the transformer model to recognize similar patterns, we wanted the input to be as uniform as possible. Ensuring this begins at file level. First of all, we don't accept any time or tempo changes in the song (line 6). Any such change could appear in our selection for a hook, thus making it inconsistent and harder for the transformer to deal with. And secondly, we only accept time signatures 2/4 (lines 9, 12) and 4/4 (lines 9, 10). However, we treat 2/4 time signatures like 4/4 time signatures and modify the meta data accordingly (line 13). Apart from accentuation this isn't a problem, since two bars in 2/4 have the same length as one bar in 4/4; all we do here implicitly is to merge every two bars of a 2/4 song into one. This simplification allows us to use more data, while preserving its uniformity. From now on we will treat all time signatures as 4/4 without further clarification.

2.3 The Hook Collection Algorithm

2.3.1 Algorithm Overview

Now, how do we now extract actual hooks from the dataset? First of all, each track in a prefiltered song disregarding certain obvious exceptions (drums, bass) has the potential to contain valuable melodic information. For example, in a rock song including a guitar solo, just extracting the main refrain melody as a hook and ignoring the solo would be a huge mistake. Even chord accompaniment can be interesting, especially considering arpeggios and other accompaniment rhythms. So we need to take a look at every track and determine with certain criteria, whether it contains a hook and where to find it. For this purpose we propose the following algorithm:

1. *If drum track, skip.*

   Drum tracks don't contain any melodic information, since drum sounds (Kick, Snare etc.) are usually mapped to arbitrary MIDI notes they have nothing to do with. Keeping any part of these would pollute the training data.

2. *Transpose to C major/A minor.*

   Notes have a different melodic function when played in different keys. For example: an F played in C major is the fourth note of the major scale, while an F in F major is the root note (first note of the scale). We don't want our transformer have to deal with these kind of differences, one note should have a clear melodic function. To achieve this, we transpose all melody tracks of major files into C major and minor files into A minor. We skip tracks with other modes.

3. *Make monophonic.*

   Harmonies are nice, however for the main melody of songs, they aren't that relevant. Since we want the transformer to focus on these melodies (hooks), we only want one note to play at once. In chords, we take the highest note. Choosing a suitable chord progression is then the composer's job.

4. *If bass track, skip.*

   Bass tracks aren't as bad as drum tracks regarding melodic information, however they often just hold the root note of the chord being played currently and don't have any significant melody of their own. There are of course exceptions to this, but they are rare. That's why we try to reduce the number of bass hooks by skipping tracks that contain a note lower than F2 after steps 2 and 3. It is important that these steps come first, step 2 so certain keys aren't favoured and step 3 so polyphonic tracks with chord accompaniment (including bass notes) aren't put at a disadvantage.

5. *Cut out 8 bars after first note played.*

   This point is for sure the most controversial. This is the greedy solution, being most runtime efficient. In application it worked surprisingly well, since a lot of songs tease their hooks in the intro. Otherwise, hooks of instruments that come in later are still captured this way. Other solutions are discussed in the section "Potential Improvements".

6. *If required note density, continue.*

   Cutting out 8 bars after the first note is played can be risky, one could only end up with a few notes, or only notes in the first 4 bars for instance. To counteract this, we introduce the following note density criteria for our hooks: Firstly, at least 12 notes have to be played in the 8 bar section. And secondly, a note has to start in 6 out of 8 bars. If one of these criteria isn't met, we skip the track.

7. *Convert to 120 bpm and collect.*

   We want note lengths and rhythm to be comparable for the transformer. That's why we convert each track to 120 bpm before collecting.

In the following sections, we will take an in-depth look at some implementations for certain steps of this algorithm. We purposefully left out the drum and bass handling because these are trivial to implement. Except for the key analysis and interval detection we used the pretty_midi library for handling MIDI in Python. All the code for these sections can be found in *Transformer_Code/CUPY/models/utils/collect_mem_2.py*.

2.3.2 Transposition To C/Am

A key part of our algorithm is to transpose major melodies into C major and minor melodies into A minor, which makes it easier for the transformer to comprehend.

```python
def detect_transpose_interval(music21_stream):
    key_signature = music21_stream.analyze('key')
    if key_signature.mode == "major":
        target_key = key.Key("C")
    elif key_signature.mode == "minor":
        target_key = key.Key("A")
    else:
        return None  # Skip non-major/minor modes
    transpose_interval = interval.Interval(key_signature.tonic,
        target_key.tonic)
    if transpose_interval.semitones > 6: #the closer way to C/Am is by
        using the reversed complement of the interval
        transpose_interval = transpose_interval.complement.reverse()
    if transpose_interval.semitones < -6: #with descending intervals
        the complement is always automatically reversed
        transpose_interval = transpose_interval.complement
    return transpose_interval.semitones
```

The first thing we do is to detect the transpose interval (line 1). For this task we use the music21 library, a library for symbolic music analysis. First we analyze the key signature with music21 (line 2). If the the melody is in major our target key is C major (line 3, 4) and if it is in minor our target key is A minor (line 5, 6). The reason that this works is because A minor is the parallel minor key of C major, meaning their scales contain the same notes (in this case only white keys), just in different order. If the melody has a different mode than major or minor, we skip it (line 7, 8), since it is uncommon for Western music and would only confuse the transformer. We then use the in-built Interval class to determine a possible transpose interval (line 9). However this isn't always the optimal solution, and often the melody was being transposed at a unnecessarily high interval where $|s| > 6$, s the number of semitones transposed. There is always a way to transpose to another key with $s \in [-6, 6]$ (Figure 1), in lines 10 to 13 we address this issue. Then we return the optimal ($|s|$ minimal) semitone value for transposing into the target key.

Figure 1: The Transposition Problem, optimal solution green

```python
def transpose_midi(pretty_midi_stream, interval):
    # Transpose the MIDI using pretty_midi
    transposed_midi = pretty_midi.PrettyMIDI()

    # Loop over all instruments and transpose their notes
    for instrument in pretty_midi_stream.instruments:
        transposed_instrument = pretty_midi.Instrument(instrument.
            program, is_drum=instrument.is_drum)
        for note in instrument.notes:
            transposed_note = pretty_midi.Note(
                velocity=note.velocity,
                pitch=note.pitch + interval,
                start=note.start,
                end=note.end
            )
            transposed_instrument.notes.append(transposed_note)
        transposed_midi.instruments.append(transposed_instrument)

    return transposed_midi
```

Now, once we know our transpose interval, we can now perform the actual transposition in pretty_midi, with the transpose interval as an argument (line 1). This is fairly straight forward: First we create a target file that we will return (line 4). Then we loop through all instruments of the old file (drums are already filtered out), create a new instrument and add the same notes with the semitones to transpose added onto their pitch (lines 8-16). This instrument we then append to the target file (line 17).

2.3.3 Making Monophonic

Making our training input monophonic is essential for hook generation. Only this way, the transformer can focus on the core melody and not get distracted by harmonies.

Figure 2: Different Approaches For Making Monophonic

To achieve this task, we need to make the following differentiation: For one, if a chord is being played, we want to take the top note. If a melody is being played, it could potentially overlap with other notes. To fix this, we want to cut off the old note when the new note starts. When to do which of these is decided by a time tolerance parameter (line 10), in our case 0.01 seconds. This is illustrated in Figure 2.

```
1   def make_monophonic(pretty_midi_stream):
2       for instrument in pretty_midi_stream.instruments:
3           # Sort notes by start time
4           instrument.notes.sort(key=lambda note: note.start)
5           # Prepare a new list to store monophonic notes
6           monophonic_notes = []
7           # Initialize a list to track simultaneous notes
8           chord_group = []
9           # Tolerance for identifying simultaneous notes (in seconds)
10          time_tolerance = 0.01
11
12          # Process notes
13          for note in instrument.notes:
14              # If the chord group is empty, add the first note
15              if not chord_group:
16                  chord_group.append(note)
17              else:
18                  # If the current note starts within the time tolerance
19                  #     of the first note in the group, add to chord
20                  if abs(note.start - chord_group[0].start) <= \
21                      time_tolerance:
22                      chord_group.append(note)
23                  else:
24                      # Process the chord group to keep only the highest
25                      #     note
26                      top_note = max(chord_group, key=lambda n: n.pitch)
27                      # If overlap occurs, cut the previous note's end
```

```
                          time
25              if monophonic_notes and monophonic_notes [-1].end >
                    top_note.start:
26                   monophonic_notes [-1].end = top_note.start
27              # Add the highest note to the monophonic list
28              monophonic_notes.append(top_note)
29              # Start a new chord group with the current note
30              chord_group = [note]
31
32          # Handle the last chord group
33          if chord_group:
34              top_note = max(chord_group, key=lambda n: n.pitch)
35              # Adjust overlap for the last group
36              if monophonic_notes and monophonic_notes [-1].end > top_note
                  .start:
37                   monophonic_notes [-1].end = top_note.start
38              monophonic_notes.append(top_note)
39
40          # Replace the instrument's notes with the monophonic notes
41          instrument.notes = monophonic_notes
```

All in all, the function loops through all tracks and performs the task separately on each
track/instrument (line 2). First we sort the notes by their start time (line 4). Then we
prepare different data structures: chord_group (line 8) for notes within the time tolerance,
monophonic_notes (line 6) for the end result. Then the actual processing starts. We go
through each note in the current track (line 13) and do the following: If there is no note
in chord_group, we add the current note (line 16). Otherwise we check whether the
current note is within the time tolerance of the first note in the chord (line 19). If yes, we
append the note to the chord group (line 20). If no, we take the highest note and append
it the our result (line 28), fixing any occuring overlaps (line 26). After that, we start a
new chord group with the current note, since it didn't belong to the previous chord.
After looping through all notes, we still might have notes in chord_group. This we handle
in lines 32-38, analogously. In the end, we replace the tracks notes with our monophonic
result (line 41).

2.3.4   The Collection Method

```
1  def greedy_collect(pretty_midi_stream, tempo, output_path_pref):
2      tracks_coll = 0
3      tracks_skipped = 0
4      for i, instr in enumerate(pretty_midi_stream.instruments):
5          if instr.notes: #this might be this case with removed bass
                tracks f.e.
6              midi = pretty_midi.PrettyMIDI(initial_tempo=120) #create a
                    new midi file for writing
7              seconds_per_beat = 60.0 / tempo
8              seconds_per_bar = 4 * seconds_per_beat #4 beats in one bar
                    (4/4, we treat 2/4 just like 4/4 here)
9              bar_duration = seconds_per_bar * 8  # Duration of 8 bars in
                    seconds
10
11             first_note_time = min(note.start for note in instr.notes)
```

```
12
13              cutoff_time = first_note_time + bar_duration #cutoff point
                    8 bars after first note
14
15              new_instrument = pretty_midi.Instrument(program=instr.
                    program, is_drum=instr.is_drum)
16              new_instrument.notes = [note for note in instr.notes if
                    first_note_time <= note.start < cutoff_time] #notes in 8
                     bars from first note
17
18              bars = [[] for _ in range(8)]
19
20              for note in new_instrument.notes:
21                  # Find the bar index (0-based index for 8 bars)
22                  bar_index = int((note.start - first_note_time) //
                        seconds_per_bar)
23                  if 0 <= bar_index < 8:  # Ensure it's within the 8 bars
24                      bars[bar_index].append(note)
25
26              bars_with_notes = sum(1 for bar in bars if len(bar) > 0)
27
28              if bars_with_notes >= 6: #something must be played in at
                    least 6 out of 8 bars
29                  for note in new_instrument.notes:
30                      note.start = max(0, note.start - first_note_time)
31                      note.end = max(0, min(note.end - first_note_time,
                            cutoff_time - first_note_time))
32                      note.start *= tempo/120 #Stretch to 120 bpm from
                            original tempo
33                      note.end *= tempo/120
34                  if len(new_instrument.notes) >= 12: #there must be at
                        least 12 notes in file
35                      midi.instruments.append(new_instrument)
36                      op_new = str(output_path_pref) + f"_track{i}.mid"
37                      midi.write(str(op_new))
38                      #print(f"8 bar excerpt saved to: {op_new}")
39                      tracks_coll += 1
40                  else: tracks_skipped += 1
41              else: tracks_skipped += 1
42
43      return tracks_coll, tracks_skipped
```

This function handles steps 5-7 of the algorithm. After initializing counters (lines 3, 4) , we iterate through all remaining tracks/instruments that still have notes (lines 5, 6). For each of them, we create a target file in 120 bpm (line 7). We calculate the seconds pear beat and bar considering the original tempo (lines 8, 9), which was given to the function as a parameter. From these results we then calculate the length of 8 bars (line 10). After that, we compute the time when the first note starts and the time where the 8 bar phrase should end (lines 12, 14). We then create an instrument with the notes from the 8 bars after the first note is played (lines 16, 17). Then we create a list of 8 lists for each bar, which we fill with the corresponding notes by using integer division (lines 19-25). After that, we calculate how many bars have notes (line 27). This we use to check one of our note density criteria (line 29). If this criterion is met, we shift the note start and end

times so that the first note starts at the beginning of the sequence and not somewhere else, respecting all edge cases (lines 31, 32). After doing that we stretch all note start and end times by a factor of $\frac{\text{tempo}}{120}$ to actually be in 120 bpm (lines 33, 34). This is necessary because we imported the notes in to a 120 bpm track, even though there were still in the original tempo of the song. Next we check our second note density criterion (line 35). If it is met, we append the created instrument to a MIDI file and write it (lines 36-40). If either criterion isn't met, we skip the track/instrument (lines 41, 42). In the end we return some counters for statistics (line 44).

2.4 Evaluation

2.4.1 Collection Statistics

When running our algorithm on the clean MIDI subset, we received the following statistics:


Figure 3: File Processing Statistics

The clean MIDI subset contains 34518 files, 17526 of these were able to be processed without an error. In total, 6846 out of these 17526 files (39.7%) were skipped due to time signature or tempo. No files were skipped due to a non-major/minor mode.

Figure 4: Track Processing Statistics

In total, 98440 tracks were processed. We collected hooks from 37875 (38%) of these. 16651 (16.9%) drum tracks were skipped. 19392 (19.7%) suspected bass tracks were skipped. 25012 (25.4%) were skipped because the first 8 bars after the first played note didn't fulfill our note density criteria.

2.4.2 Positive/Negative Examples

In this section we will list some positive and negative results from our Hook Collection. They can be found in *Transformer_Code/paper_data*. As a test set we used ABBA songs. The corresponding MIDI files can be found in ...

Positive Examples:

- In "Dance (While the Music Still Goes on)1_track0" the vocal hook from the beginning is captured well.

- In "Mamma Mia_track0" the piano hook from the intro is captured perfectly.

- In "Name of the Game.1_track4" a nice accompanying string melody is captured.

- In "Name of the Game.1_track8" a interesting accompaniment arppegio is captured.

- In "Summer Night City_track4" a minimalistic, but good hook is captured.

Negative Examples:

- In "Dance (While the Music Still Goes on).1_track6" the top note from a chord accompaniment is captured. It sounds fine, but isn't a great hook

- In "I Do, I Do, I Do, I Do_track1" once again top notes from a chord accompaniment are captured, that don't sound like a hook on their own.

- In "Just Like That (Full Sax version) (1983)_track1" a rhythmic accompaniment that sounds monotonic and bland is captured.

- In "Opus 10_track1" the rhythm and notes seem a little random, not catchy at all.

- In "The Living Daylights_track16" a interesting rhythm is captured. However, the pattern mostly alternates between two notes, which is a bit boring.

2.4.3   Potential Improvements

On file level, filtering out files with alternating time signatures and tempo and focusing on 4/4 time signature is definitely an important step toward uniformity. On track level drum track skipping, tranposing and making monophonic also worked very well. However, the next steps of our algorithm can be questioned:

First of all, skipping tracks with notes lower than F2 after steps 2 and 3 gets rid of more tracks than it actually should; as can be seen in "Collection Statistics", about every fifth track is classified as a bass track, even though every file usually has 6 or more tracks, only one of them being a bass track. This issue could be solved by training a classifier (Binary Logistic Regression, Neural Network or Transformer) to detect whether a track is a bass track or not. Another possibly useful feature is that bass tracks are always monophonic. Tracks that aren't monophonic (before making everything monophonic) are therfore definitely never a bass track and wouldn't have to be skipped.

Secondly, the approach of cutting out 8 bars after the first played note in a so far accepted track isn't the best solution. It ignores the fact that a better hook could theoretically could be in the middle of the track. Another approach would be slicing the input into 8 bar pieces after the first played note and checking certain note density criteria for each. This would have the advantage that more data is collected and the disadvantage that many 8 bar sequences wouldn't be hooks, since songs also have a lot of other melodies. Another possibility without this disadvantage would be training a neural network to predict the 8 bars most likely to be a hook in a MIDI sequence and then collecting that.

And thirdly, our note density criteria are rather strict. It can be argued that a note played in 4 or 5 out of 8 bars would suffice. However, at the with those settings we generated a lot of pauses and this was an easy and effective way for us to counteract this.

3   Tokenization

3.1   The REMI Tokenizer

In order to tokenize our MIDI events we use the REMI tokenizer, implemented in the MidiTok (Fradet et al. 2021) library, which was first introduced in the Pop Music Transformer paper written by Y.-S. Huang and Yang 2020. The REMI tokenizer provides a structured approach to representing musical events in MIDI format by breaking down the sequence of notes into a series of tokens. This tokenization method captures essential

musical features such as Bar, Pitch, Tempo, Position, Velocity and Note Duration. Therefore it is well fit for our task.


Figure 5: REMI tokenization example

3.2 Important Configuration Parameters

```
TOKENIZER_PARAMS = {
    "pitch_range": (21, 109),
    "beat_res": {(0, 32): 8},
    "num_velocities": 1,
    "special_tokens": ["PAD", "BOS", "EOS"],
    "use_note_duration_programs": [-1, 0, 1, 2, 3, 4, 5, 6, 7, 8,
        9, 10, 11, 12, 13, 14, 15, 16, 17, 18, 19, 20, 21, 22, 23,
        24, 25, 26, 27, 28, 29, 30, 31, 32, 33, 34, 35, 36, 37, 38,
        39, 40, 41, 42, 43, 44, 45, 46, 47, 48, 49, 50, 51, 52, 53,
        54, 55, 56, 57, 58, 59, 60, 61, 62, 63, 64, 65, 66, 67, 68,
        69, 70, 71, 72, 73, 74, 75, 76, 77, 78, 79, 80, 81, 82, 83,
        84, 85, 86, 87, 88, 89, 90, 91, 92, 93, 94, 95, 96, 97, 98,
        99, 100, 101, 102, 103, 104, 105, 106, 107, 108, 109, 110,
        111, 112, 113, 114, 115, 116, 117, 118, 119, 120, 121, 122,
        123, 124, 125, 126, 127],
    "use_chords": False,
    "use_rests": True,
    "beat_res_rest": {(0, 32): 8},
    "use_tempos": False,
    "use_time_signatures": False,
    "time_signature_range": {4: [4]},
    "use_programs": False,
    "use_pitchdrum_tokens": False,
    "default_note_duration" : 0.5,
    "num_tempos": 32,  # number of tempo bins
    "tempo_range": (120, 120),  # (min, max)
    "encode_ids_split": 'no',
    "ac_note_duration_track": True,
    "ac_note_density_track": True,
    "ac_repetition_track": True,
    "ac_repetition_track_num_bins": 8
}
```

These are the configuration parameters we used to set up our tokenizer. They can be found in *Transformer_Code/CUPY/models/utils/generate_tokenizer.py*. Many are default settings, but some of them we had to adapt to our task:

```
"beat_res": {(0, 32): 8}
```

This setting enables the tokenizer to capture notes up to 32nd notes in the first 8 bars (these are equal to 32 beats, our inputs are never longer than this). Notes in between are

rounded to the nearest 32nd note. We chose 32nd notes because in popular music notes are rarely shorter than this.

```
"num_velocities": 1
```

We don't care about different velocities in melodies, since these have nothing to do with the melody played. That's why we downsample everything to have the same velocity.

```
"use_chords": False
```

Since our input is monophonic, there are no chords, so we don't want the tokenizer to try and encode them.

```
"use_rests": True,
"beat_res_rest": {(0, 32): 8}
```

Rests are important to capture the rhythm of the hooks correctly. That's why we capture them at the same beat resolution as notes.

```
"use_tempos": False,
"tempo_range": (120, 120),
"use_time_signatures": False,
"time_signature_range": {4: [4]}
```

These settings align with how we collected our data. We can treat everything as a 120 bpm, 4/4 time signature hook. Any different time signatures or tempos aren't needed.

```
"encode_ids_split": 'no'
```

The default setting here was 'bar', meaning that tokens weren't able to go beyond bars. This is problematic, since important melodic structures and repetitions of hooks do go beyond bars, that's why we set this parameter to 'no'.

```
"ac_note_duration_track": True,
"ac_note_density_track": True,
"ac_repetition_track": True,
"ac_repetition_track_num_bins": 8
```

These parameters allow the tokenizer to use extra tokens to track certain behaviours in our hooks like note duration (important for capturing rhythm correctly), repetition (very important in hooks due to the repetitiveness) and note density development.

3.3   Vocabulary Size

In monophonic melodies, only one note is active at a time, so the vocabulary doesn't need to account for simultaneous note combinations (e.g., chords). We opted for a vocabulary size of 4096 which might be very generous given that out data consists of single-note melodies. In addition to that we expect short sequences: An 8-bar melody is

relatively short, especially in monophonic form. For example, in 4/4 time at 120 BPM, an 8-bar melody lasts 16 seconds.

3.4 Data Augmentation

In order to enlarge our training data we choose to augment the midi tracks by shifting the melody one octave up, two octaves up, one octave down and two octaves down. Fortunately the MidiTok library already built a data augmentation function which we used in our project.

3.5 Token Statistics

In this section we summarise the tokenization procedure by showing distributional information.

Firstly we observed that the top 20 tokens are all tokens which encode Tempo, Positional or Bar events. REMI insert these tokens at regular intervals (e.g., every bar or every beat) to ensure the timing information is preserved. We opted for capturing a 32nd note resolution and therefore many these tokens are used in order to capture this detailed beat resolution. Figure (7) depicts the total token distribution. The top 20 tokens are responsible for the spike at around 350. Starting from token 696 we observed token ids which actually encode Note events. These are responsible for the second peak which flattens out toward the maximum vocabulary size. It is important to note that the Note Events do also encode time, duration, pitch and velocity. In total the vocabulary has 2955 tokens for Note events and 1141 tokens for Tempo, Positional or Bar events. With regard to the sequence length distribution we see that 50% of the input files have a sequence length smaller than 103. Nevertheless the variance of sequence lengths is quite high as seen in Figure (10) and (9). We used this information to choose the models context size which is discussed in section 5.1.


Figure 6: Top 20 tokens

15

Figure 7: Token distribution


Figure 8: Sequence Length Box Plot

Figure 9: Cumulative Distribution of Sequence Lengths

Figure 10: Log Distribution of Sequence Lenghts

4 Relative Attention

4.1 Overview

Note: As an abbreviation for the matrix dimensions we will be using B for the batch size, H for the number of heads, T for the sequence length and D for the head depth.

Since we are dealing with music in this task, the Absolute Attention mechanism commonly used in transformers (Vaswani et al. 2017) doesn't quite cut it. This mechanism only captures the relationship of tokens at absolute positions with one another. However, we also need to capture general relationships between notes. For example: How do adjacent notes/tokens correlate in general, can we infer anything from our data?

A solution to this problem was first introduced by Shaw, Uszkoreit, and Vaswani (2018) and then refined by C. A. Huang et al. (2018). The main objective is to capture these relationships in the form of learnable relative positional embeddings E (of dimensions $1 \times H \times T \times D$) in every attention block. Each row of such a matrix corresponds to a certain key-query distance, as can be seen in Figure 11.

Figure 11: Relative Attention With Skewing, C. A. Huang et al. (2018)

The first step of the attention process is to multiply Q by E^T, resulting in a matrix of dimensions $B \times H \times T \times T$ (E is broadcasted over B beforehand). Its (i_q, r) entry contains the dot product of the query in position i_q with the embedding of relative distance r. However, each relative logit (i_q, j_k) in the matrix S_{rel} should be the dot product of the query in position i_q and the corresponding key j_k with relative distance r, to match up with the indexing in QK^T. C. A. Huang et al. (2018) propose the "skewing" algorithm to move the relative logits of QE^T to their correct positions, as illustrated in Figure 11:

1. Pad a dummy column vector before the leftmost column (shape: $B, H, T, T+1$).

2. Reshape the matrix to have shape $(B, H, T+1, T)$.

3. Discard the top row of the third dimension (shape: B, H, T, T).

The result of this process is the relative attention score matrix S_{rel} with correct indexing, as outlined in Equation (1). This matrix is now added to the conventional attention scores QK^T. This sum then goes through the normal attention process from Vaswani et al. (2017), as outlined in Eq. 2.

$$S_{\text{rel}} = \text{skew}(QE^T) \tag{1}$$

$$\text{RelativeAttention} = \text{softmax}\left(\frac{QK^T + S_{\text{rel}}}{\sqrt{d_k}}\right)V \tag{2}$$

4.2 Implementation

4.2.1 Forward Pass

In this section, we will go through all relevant differences to normal Absolute Attention in the Forward Pass, meaning we won't be showing any computations that are exactly the same as in Absolute Attention. All code for this and the next subsection can be found in *Transformer_Code/CUPY/models/GoePT/layers.py*

```
self.x1 = q @ e.transpose(0, 1, 3, 2) #Q * E_T
```

Here we compute $X_1 = QE^T$ as portrayed in the section "Overview", the broadcasting of E happens automatically here.

```
self.x2 = cp.pad(self.x1, ((0, 0), (0, 0), (0, 0), (1, 0)), mode='
    constant', constant_values=0) #padding, adding a column of zeros
    to the left of the T x T matrix, results in (B, nh, T, T + 1)
```

Step 1 of the Skewing Algorithm: $X_2 = \text{pad\_left\_column}(X_1)$.

```
self.x3 = self.x2.reshape(B, self.n_heads, T + 1, T) #reshape to (B
    , nh, T + 1, T) matrix
```

Step 2 of the Skewing Algorithm: $X_3 = \text{reshape } X_2 \text{ to } (B, H, T + 1, T)$.

```
self.s_rel = self.x3[:, :, 1:, :] #discard first row of 3rd dim,
    new dim: (B, nh, T, T)
```

Step 3 of the Skewing Algorithm: $S_{\text{rel}} = \text{discard\_top\_row}(X_3)$

```
attn = (q @ k.transpose(0, 1, 3, 2) + self.s_rel)*(1.0/math.sqrt(k.
    shape[-1])) #relative attention
```

Here we implement the mask input from Equation 2 from the section "Overview". From here on everything works exactly the same as in Absolute Attention.

4.2.2 Backward Pass

Note: We will be using $M = \frac{QK^T + S_{rel}}{\sqrt{d_k}}$ as an abbreviation for the mask input from Equation (2). In addition, we will be using E_b as an abbreviation for E broadcasted to all batches. We also implicitly assume that transposition and matrix multiplication only happen over the last two dimensions of tensors. Therefore we will be referring to the the 3rd dimension as rows and 4th dimension as columns.

In this section, we will go through all relevant differences to Absolute Attention in the Backward Pass, so we won't be showing gradients that are exactly the same. $\frac{\partial L}{\partial M}$ is already given to us, since the attention process after computing M is identical to the process in absolute attention (first mask, then softmax etc.).

```
grad_s_rel = (1.0/math.sqrt(self.k.shape[-1])) * grad_mask
```

$$\frac{\partial L}{\partial S_{\text{rel}}} = \frac{\partial L}{\partial M} \cdot \frac{\partial M}{\partial S_{\text{rel}}} \tag{3}$$

$$= \frac{\partial L}{\partial M} \cdot \frac{1}{\sqrt{d_k}} \tag{4}$$

The gradient of L with respect to S_{rel} can be easily computed with the chain rule (Eq. 3). The gradient of M with respect to S_{rel} is simply the constant $\frac{1}{\sqrt{d_k}}$ (Eq. 4).

19

```
1    grad_x3 = cp.pad(grad_s_rel, ((0, 0), (0, 0), (1, 0), (0, 0)), mode
         ='constant', constant_values=0) #pad top row
```

$$\frac{\partial L}{\partial X_3} = \text{pad\_top\_row}\left(\frac{\partial L}{\partial S_{\text{rel}}}\right) \tag{5}$$

The gradient of L with respect to X_3 can be received by performing the reverse operation of step 3 of the Skewing Algorithm. There we discarded the top row, now we pad the top row of our current gradient $\frac{\partial L}{\partial S_{\text{rel}}}$ with 0's (Eq. 5). This results in a tensor of shape $(B, H, T + 1, T)$.

```
1    grad_x2 = grad_x3.reshape(B, self.n_heads, T, T + 1) #reverse
         reshaping
```

$$\frac{\partial L}{\partial X_2} = \text{reshape } \frac{\partial L}{\partial X_3} \text{ to } (B, H, T, T + 1) \tag{6}$$

The gradient of L with respect to X_2 can be received by performing the reverse operation of step 2 of the Skewing Algorithm. There we reshaped the tensor X_2 to $(B, H, T + 1, T)$, now we reshape the gradient with respect to X_3 to the shape $(B, H, T, T + 1)$ (Eq. 6).

```
1    grad_x1 = grad_x2[:, :, :, 1:] #discard the left column
```

$$\frac{\partial L}{\partial X_1} = \text{discard\_left\_column}\left(\frac{\partial L}{\partial X_2}\right) \tag{7}$$

The gradient of L with respect to X_1 can be received by performing the reverse operation of step 1 of the Skewing Algorithm. There we padded the tensor X_1 with 0's on its left column, now we discard the current gradient's $\frac{\partial L}{\partial X_2}$ left column (Eq. 7). This results in a tensor of shape (B, H, T, T).

```
1    grad_q = (1.0/math.sqrt(self.k.shape[-1])) * (grad_mask @ self.k)
2    broadcasted_emb = cp.broadcast_to(self.rel_pos_emb, (B, self.
         n_heads, T, self.depth))
3    grad_q2 = grad_x1 @ broadcasted_emb
4    grad_q += grad_q2
```

$$\frac{\partial L}{\partial Q} = \frac{\partial L}{\partial X_1} E_b + \frac{1}{\sqrt{d_k}} \cdot \frac{\partial L}{\partial M} K \tag{8}$$

The right summand of this gradient is identical to $\frac{\partial L}{\partial Q}$ in Absolute Attention (line 1), that's why there is no need to explain it any further. With Relative Attention however, there is another path (over S_{rel}) to derive this gradient that also needs to be taken into account. We know that $X_1 = Q(E_b)^T$, according the the matrix multiplication rule the derivative with respect to Q on this path is equal to the matrix product of the derivative with respect to X_1 and E_b (line 3). This summand is now added to receive the complete formula (line 4, Eq. 8). Both summands and the end result have the shape (B, H, T, D).

```
1    grad_rpe_all_b_trans = self.q.transpose(0, 1, 3, 2) @ grad_x1
2    grad_rpe_all_b = grad_rpe_all_b_trans.transpose(0, 1, 3, 2)
3    self.grad_rpe = cp.sum(grad_rpe_all_b, axis=0, keepdims=True) #Sum
         over all batches to get desired dim: (1, nhs, T, d)
```

$$\frac{\partial L}{\partial (E_b)^T} = Q^T \frac{\partial L}{\partial X_1} \tag{9}$$

$$\frac{\partial L}{\partial E_b} = \left(Q^T \frac{\partial L}{\partial X_1} \right)^T \tag{10}$$

$$\frac{\partial L}{\partial E} = \sum_B \left(Q^T \frac{\partial L}{\partial X_1} \right)^T \tag{11}$$

In this section, we will step by step derive $\frac{\partial L}{\partial E}$. We already know that $X_1 = Q(E_b)^T$, from previous sections. This time we derive with respect to $(E_b)^T$. Using the matrix multiplication rule we receive the result from Eq. 9 with shape (B, H, D, T) (line 1). Transposing the result, we get the gradient with respect to E_b with shape (B, H, T, D) (Eq. 10, line 2). To reach the final result, we must now contract the tensor over the batch dimension (Eq. 11, line 3), since we indirectly broadcasted it in the forward pass. The resulting tensor now has the correct shape $(1, H, T, D)$.

5 Training And Inference

5.1 Training

To finally train our model we used the following procedure:

1) Pre-process the MIDI-dataset.

2) Train the tokenizer via Byte-Pair-Encoding(BPE).

3) Generate a Train-, Validation and Testset and perform Data augmentation.

4) Tokenize the datasets using the pre-trained REMI-tokenizer.

5) Save the tokenized datasets to disk.

6) Run the training routine which is based on a nano-GPT implementation.

Initially, we organized the data into a matrix structure, where each row corresponds to a MIDI file and the columns represents a sequence of tokens. The column size was bounded by the model's context size.

$$M = \begin{bmatrix} x_{1,1} & x_{1,2} & x_{1,3} & \dots & x_{1,T} \\ x_{2,1} & x_{2,2} & x_{2,3} & \dots & x_{2,T} \\ x_{3,1} & x_{3,2} & x_{3,3} & \dots & x_{3,T} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ x_{N,1} & x_{N,2} & x_{N,3} & \dots & x_{N,T} \end{bmatrix}$$

As it turned out this initial approach had a major problem. The sequence length of the tokenized midi files is not fixed. Every file could have a different sequence length. Either greater or smaller than the models context size. Due to that we either needed to truncate many sequences or add padding tokens based on the context size of the model. Choosing a context size that is very large resulted in to many padding tokens which dominated the data. Making the context size small resulted in to many truncations such that we weren't able to fully supply the 8 bar hooks to the model in many cases. Using this approach and especially smaller context sizes for training yielded excellent loss curves and just after a few epochs we observed loss values of 0.1 to 0.001. In hindsight, this is not really surprising since padding tokens dominated our dataset and the model could just predict padding tokens. With regard to hook generation this excellent loss values were completely useless since the model only produced padding tokens during inference.

In order to solve this problem we changed the data structure from a matrix to a flat array. This approach has two advantages. We are not constrained by the models context size and we do not need to include too many padding tokens. Now our input data looks like this (example data)

$$
M = \begin{bmatrix}
0, & 0, & 0, & 0, & 0, & 0, & 0, & 0, & 0, & 0 & 1, \\
629, & 961, & 341, & 670, & 1702, & 35, & 1025, & 706, & 1702, & 2384, & 34 \\
847, & 3104, & 644, & 357, & 1073, & 1176, & 1616, & 1495, & 629, & 1757, & 1106 \\
670, & 1702, & 2384, & 1616, & 877, & 349, & 1345, & 875, & 2384, & 34, & 847 \\
640, & 857, & 644, & 357, & 1073, & 1176, & 34, & 847, & 1199, & 2, & 0 \\
0, & 0, & 0, & 0, & 0, & 0, & 0, & 0, & 0, & 0, & 0 \\
0, & 0, & 0, & 0, & 0, & 0, & 1, & 629, & 751, & 1490, & 667 \\
357, & 3213, & 674, & 105, & 1267, & 751, & 92, & 3993, & 3213, & 674, & 105 \\
1267, & 759, & 341, & 29, & 1186, & 635, & 2857, & 2867, & 2378, & 1733, & 1321 \\
1117, & 341, & 883, & 2824, & 347, & 994, & 2663, & 349, & 1148, & 2505, & 2961 \\
629, & 937, & 104, & 2, & 0, & 0, & 0, & 0, & 0, & 0, & 0
\end{bmatrix}
$$

where 1 represents the start token, 2 represents the end token and 0 represents the padding token. We randomly add padding tokens before each sequence in order to improve the inference procedure. The training routine then loads a random subset of this array of the size of the context size. Figure (12) and (13) depicts a few of our training runs. In total we trained the models for two days using different parameter configurations. As shown we managed to achieve a stable training process without indications of overfitting.

Figure 12: Validation Loss

Figure 13: Training Loss

Our final model is summarized in table (5.1).

Parameter	
Model	loyal-chocolate
Dropout	0.35
intial learning rate	0.00005
Heads	8
Embedding Size	256
Layers	4
Context Size	256
Vocabulary Size	4096

In general we observed that increasing the number of layers, heads or the embedding size did not yield a lower loss. The context size was chosen given the information we obtained from our token statistics (3.5). It is important to note again that the way we feed the data into the model does not rely heavily on the most accurate context size given the sequence length distribution of the data. By randomly selecting data from our input array chances are high that we input two incomplete sequences at once. This is not a problem since the model learns to distinguish between sequences due to the start and end token.

5.2 Inference Process

Our inference pipeline generates MIDI sequences based on a given input MIDI file using a trained transformer model. The key steps involved in the process are:

5.2.1 Model Loading

We load a pre-trained model from a checkpoint file along with the corresponding vocabulary file.

Listing 1: Loading the pre-trained model and tokenizer

```
def load_model(checkpoint_path, vocab_file, batch_size):
    with open(checkpoint_path, mode='r', encoding='utf-8') as weights:
        state_dict = json.load(weights)

    model = GoePT.from_state_dict(state_dict, batch_size=batch_size)

    tokenizer = config.tokenizer_name(params=vocab_file)
    return model, tokenizer
```

Code Location: *Transformer_Code/CUPY/models/GoePT/Inference.py*

5.2.2 Tokenization

The input MIDI files are tokenized using the REMI tokenizer, converting musical events into a sequence of tokens. Special tokens such as Start-of-Sequence (SOS) and End-of-Sequence (EOS) are appended as needed.

5.2.3 Sequence Generation

Given a tokenized input sequence, the model generates new tokens autoregressively. The following techniques are applied to enhance the quality of generated sequences:

- **Softmax with Temperature:** Adjusts the sharpness of the probability distribution to control randomness.

- **Top-p Sampling:** Selects tokens from the smallest subset of the vocabulary whose cumulative probability exceeds p, ensuring a balance between diversity and coherence. This is explained in more detail in section 5.3

The model generates tokens iteratively, stopping when the End-of-Sequence (EOS) token is encountered or the maximum amount of tokens is reached.

Listing 2: Generating sequence with EOS stopping condition

```
def generate_sequence(model, input_sequence, tokenizer, max_tokens, p,
    T):
    seq_len = model.context_length
    pad_token = tokenizer.pad_token_id
    sos_token = tokenizer.special_tokens_ids[1]
    eos_token = tokenizer.special_tokens_ids[2]

    input_sequence = [sos_token] + tokenizer(input_sequence)[0].ids
    input_sequence = cp.array(input_sequence).reshape(1, -1)

    generated_sequence = cp.copy(input_sequence)

    for idx in range(max_tokens):
        logits, _ = model.forward(input_sequence, targets=None)
        logits = cp.squeeze(logits, axis=1)  # Transform to 2D shape (
            batch, vocab)

        predictions = softmax_with_temperature(logits, temperature=T)
        next_tokens = top_p_sampling(predictions, p=p)

        if next_tokens == eos_token:
            break  # Stop generation when EOS is found

        # Append predicted token
        generated_sequence = cp.concatenate([generated_sequence,
            next_tokens], axis=1)

    return generated_sequence.get()  # Move to main memory
```

5.2.4 Post-processing

After generation, the token sequences are decoded back into MIDI format using MidiTok. Additionally, we apply a monophonic constraint to ensure that only the highest-pitched note is retained in overlapping chords. This is necessary because we chose not to quantize our MIDI input for tokenization, in order to leave the rhythm as human as possible. With notes faster than 32nd notes (our beat resolution) polyphonic tokens are unfortunately still possible.

5.3 Top-p Sampling

5.3.1 Overview

Before the introduction of Top-p sampling, maximum likelihood decoding and beam search were the standard decoding techniques for generation tasks, but, both of these decoding strategies are prone to generating texts (or melodies) that are repetitive and often end up in loops. Top-p sampling is a decoding strategy for autoregressive language model outputs proposed by Holtzman et al. (2020) that solves these issues. To explain it, we will go through the following example:


Figure 14: Top-p Sampling, Stokes (2020)

Our task is to predict the next word/token of the sentence S = "The boy went to the ...". Our model outputs this probability distribution for the next word/token:
$P(park|S) = 0.37, P(store|S) = 0.3, P(grocery|S) = 0.1, P(beach|S) = 0.06$ etc.

The sampling works like this: We set a parameter p. Then we sort the next tokens by their probability from highest to lowest. After that we choose the first k tokens of the list, whose cumulative probability exceeds the threshold p. Our prediction probability for all other tokens we set to 0, the prediction probabilities of the chosen tokens we divide by the cumulative probability, thus rescaling them to be a full probability distribution.

For our example: We choose $p = 0.75$. The first tokens in the sorted list exceeding this threshold are park, store and grocery with cumulative probability 0.77. These prediction probabilities are rescaled to $P(park|S) = \frac{0.37}{0.77}, P(store|S) = \frac{0.3}{0.77}, P(grocery|S) = \frac{0.1}{0.77}$ and all other token predicition probabibilities are set to 0.

As an additional note: Top-k sampling is a similar technique except that the sample is taken from the k-highest probability tokens regardless of their cumulative probability. The advantage of top-p sampling is that one avoids the problem of choosing the optimal value of k which can vary depending on the shape of the output distribution and the particular task and dataset.

5.3.2 Implementation

```python
def top_p_sampling(prob_matrix, p=0.2):
    batch_size, vocab_size = prob_matrix.shape
    sampled_indices = cp.zeros((batch_size,1), dtype=int)

    for i in range(batch_size):
        probs = prob_matrix[i].copy()  # Copy to avoid modifying
            original data

        # Sort tokens by descending probability and get indices
        sorted_indices = cp.argsort(probs)[::-1]
        sorted_probs = probs[sorted_indices]

        # Compute cumulative probabilities
        cumulative_probs = cp.cumsum(sorted_probs)

        # Find the cutoff where cumulative probability exceeds p
        # Use argmax to find the first index where condition is True;
            add 1 to include that index
        # If none exceed p, argmax returns 0 (False for all), cutoff
            becomes 1 (keep the first)
        cutoff = cp.argmax(cumulative_probs > p) + 1

        # Slice the top indices and probabilities up to the cutoff
        top_indices = sorted_indices[:cutoff]
        top_probs = sorted_probs[:cutoff]

        # Normalize the probabilities
        top_probs /= cp.sum(top_probs)

        # Sample from the top distribution
        sampled_index = cp.random.choice(top_indices, size=1, p=
            top_probs)
        sampled_indices[i] = sampled_index[0]

    return sampled_indices
```

This is how we implemented the sampling method for our use case (can be found in *Transformer_Code/CUPY/models/GoePT/Inference.py*). The steps strictly follow the algorithm described in "Overview", any additional explanitions are given in the comments.

5.4 Top-p and Temperature Effects on Melodic Outputs

Lower *top-p* and *temperature* were found to generate more coherent, melody-like sequences. This behavior likely arises from the inherent structure of the tokenizer's vocabulary and the probabilistic nature of the decoding process:

1. **Token Distribution Skew:**
   While individual note tokens may have high probabilities during inference, the training data contains a disproportionate number of time-related tokens (tempo, bar, time signature). We assume this skew causes these non-melodic tokens to

collectively dominate the probability mass, even if no single time token appears in the absolute top ranks.

2. **Impact of High p:** When increasing p, the sampling pool expands to include lower-ranked tokens. Due to the dataset's structural bias, note tokens become statistically "drowned out" by the sheer quantity of tempo/metadata tokens in the expanded pool. This results in degenerate sequences overpopulated with non-musical events.

3. **Effect of Lower Temperature:** Lower temperature focuses the model on the most likely tokens, which are often notes because they fit the melody best. This helps avoid distractions from less important tokens like tempo markers, keeping the output more musical.

4. **Why Lower Settings Work Better:** By reducing both p and *temperature*, the decoding process effectively filters out non-essential time tokens, skewing the model toward the smaller subset of tokens that encode pitch, rhythm, and articulation. This aligns with the simplicity of monophonic hooks. But choosing p and *temperature* to small, the generated sequence will not sound like any recognisable melody. In our specific case we used p-values of 0.1 to 0.4 and temperate-values of 0.6 to 0.8.

6 Conclusion

Due to our adaptions, our transformer manages to output good sounding melody snippets (about 2 bars). However, it didn't quite manage capture the overall structure of hooks, making it sound a little random and strange at some points. Our top 3 generations can be found at *Transformer_Code/Top3Generations*. All in all, we got a good insight into the world of Music Transformers and we were able to apply our knowledge from class and other courses.

References

Fradet, Nathan et al. (2021). "MidiTok: A Python package for MIDI file tokenization". In: *Extended Abstracts for the Late-Breaking Demo Session of the 22nd International Society for Music Information Retrieval Conference*. URL: https://archives.ismir.net/ismir2021/latebreaking/000005.pdf.

Holtzman, Ari et al. (2020). *The Curious Case of Neural Text Degeneration*. arXiv: 1904.09751 [cs.CL]. URL: https://arxiv.org/abs/1904.09751.

Huang, Cheng-Zhi Anna et al. (2018). "An Improved Relative Self-Attention Mechanism for Transformer with Application to Music Generation". In: *CoRR* abs/1809.04281. arXiv: 1809.04281. URL: http://arxiv.org/abs/1809.04281.

Huang, Yu-Siang and Yi-Hsuan Yang (2020). "Pop Music Transformer: Beat-Based Modeling and Generation of Expressive Pop Piano Compositions". In: *Proceedings of the 28th ACM International Conference on Multimedia*. MM '20. Seattle, WA, USA: Association for Computing Machinery, pp. 1180–1188. ISBN: 9781450379885. DOI: 10.1145/3394171.3413671. URL: https://doi.org/10.1145/3394171.3413671.

Lakh Clean Midi Dataset Analysis (2024). Accessed: 2025-02-27. URL: https://lakhcleananalysis.sourceforge.io/.

Raffel, Colin (2016). *Learning-Based Methods for Comparing Sequences, with Applications to Audio-to-MIDI Alignment and Matching*. URL: https://api.semanticscholar.org/CorpusID:63439223.

Shaw, Peter, Jakob Uszkoreit, and Ashish Vaswani (2018). *Self-Attention with Relative Position Representations*. arXiv: 1803.02155 [cs.CL]. URL: https://arxiv.org/abs/1803.02155.

Stokes, Jackson (2020). *A guide to language model sampling in AllenNLP*. Accessed: 2025-02-27. URL: https://medium.com/ai2-blog/a-guide-to-language-model-sampling-in-allennlp-3b1239274bc3.

Vaswani, Ashish et al. (2017). "Attention Is All You Need". In: *CoRR* abs/1706.03762. arXiv: 1706.03762. URL: http://arxiv.org/abs/1706.03762.